

ABHIJAY SINGH PANWAR

Indore, India | +91 9630544817 | abhijaysinghpanwar822@gmail.com
github.com/AbhijaySinghPanwar | linkedin.com/in/abhijay-singh-panwar

EDUCATION

Vellore Institute of Technology, Vellore

2023 – 2027

B.Tech in Information Technology — CGPA: 8.77

EXPERIENCE

Data Engineering & Analytics Intern, Tech Mahindra

May 2026 – Jul 2026

- Automated ingestion of heterogeneous datasets across 3+ file formats, cutting manual preprocessing time by roughly 30%.
- Engineered reusable SQL queries and Spark transformations supporting revenue and customer analytics across telecom datasets covering 100K+ records.
- Strengthened data quality via schema validation and missing-value handling across 5+ pipelines, cutting downstream errors by nearly 20%.
- Delivered 3 interactive Power BI dashboards tracking churn KPIs, revenue trends, and retention analysis for stakeholders.

PROJECTS

ResumeAI — AI-Powered Resume Analysis Platform

[GitHub](#) | [Live Demo](#)

Python, FastAPI, Docker, AWS EC2, Amazon RDS PostgreSQL, Nginx, SQLAlchemy, Alembic, ONNX Runtime, Google Gemini API

- Built and deployed an AI-powered ATS resume analysis platform on AWS EC2 using Docker, FastAPI, Nginx, and Amazon RDS PostgreSQL, enabling semantic resume-to-job matching across 10+ job categories.
- Developed resume parsing, ATS scoring, semantic job matching, skill-gap analysis, AI-powered resume enhancement, and interview preparation, validated using 60+ real resumes.
- Optimized inference with ONNX Runtime, reducing latency by ~40%, and secured production deployment using AWS Systems Manager Parameter Store and Alembic for configuration and database migrations.

Smritikosh — Multimodal Personal Memory Retrieval System

[GitHub](#)

Python, FastAPI, React, CLIP, BLIP-2, Whisper, Qdrant, KùzuDB

- Engineered a seven-stage retrieval pipeline over a personal photo, video, audio, and document archive, answering natural-language queries against 500+ real media items.
- Implemented multimodal extraction with CLIP, BLIP-2, and Whisper, episode segmentation, and a KùzuDB memory graph, with tiered scheduling and query-triggered deepening to cut redundant inference.
- Built entity resolution with face clustering and correction propagation, plus privacy-tiered routing that keeps sensitive content on local models.

Healthcare Appointment & Follow-up Manager

[GitHub](#) | [Live Demo](#)

Node.js, Express.js, TypeScript, PostgreSQL, Prisma, React, JWT, Google Calendar OAuth 2.0

- Eliminated double-booking at the database level using a PostgreSQL GiST exclusion constraint, backed by a hold-then-confirm state machine with a 5-minute TTL and SHA-256 idempotency keys.
- Handled concurrent booking requests with SELECT FOR UPDATE SKIP LOCKED and deterministic Google Calendar event IDs, keeping calendar sync idempotent across retries.
- Added Groq LLM symptom triage and node-cron follow-up reminders; deployed on Vercel with JWT-authenticated routes.

TECHNICAL SKILLS

Languages: Python, C++, JavaScript, TypeScript, SQL, HTML/CSS

Frameworks: FastAPI, Node.js, Express.js, React, Prisma

AI/ML: Google Gemini API, Groq, Sentence-Transformers, CLIP, BLIP-2, Whisper, ONNX Runtime, Hugging Face

Data Engineering: Apache Spark, Databricks, Pandas, Power BI

Databases: Amazon RDS PostgreSQL, PostgreSQL, MongoDB, Qdrant, KùzuDB, SQLite

Developer Tools: Git, GitHub, Docker, Nginx, AWS EC2, AWS Systems Manager Parameter Store, Vercel, Hugging Face Spaces

CERTIFICATIONS

- Claude with the Anthropic API — Anthropic (2026)
- 100 Days of Code: The Complete Python Pro Bootcamp — Udemy (Jul 2026)
- The Complete Full-Stack Web Development Bootcamp — Udemy (May 2026)
- GenAI Using IBM Watsonx — IBM Developer Skills Network (Jun 2025)
- Decode C++ with Data Structures & Algorithms — PW Skills (Jun 2024)